RTable 1. **Quantitative ablation of BNoise and ITB on the IIOF dataset.** We compare restoring VMDiff's BNoise initialization (*w/ BNoise*) with the proposed Image–Text Blending (*w/ Our ITB*). The results show that ITB consistently improves SS and all VLM-based metrics, supporting that the gain comes from the proposed image–text coupling rather than from BNoise initialization.

Methods	**SS** ↑	**CLIP-FlanT5** ↑			**Qwen2.5-VL** ↑			**LLaVA-1.5** ↑		
		P_g	P_{single}	P_{Fuse}	P_g	P_{single}	P_{Fuse}	P_g	P_{single}	P_{Fuse}
w/ BNoise	1.588	0.606	0.423	0.506	0.507	0.186	0.370	0.484	0.331	0.295
w/ Our ITB	**1.736**	**0.647**	**0.466**	**0.562**	**0.618**	**0.252**	**0.468**	**0.520**	**0.359**	**0.332**

RTable 2. **Ablation on early-readout timesteps.** Results on the *strawberry* subset, where *strawberry* is fused with the other objects in the dataset. Runtime is reported in seconds. We mark the selected timestep as 776†. The results support the use of early readout as an efficient search proxy: compared with near-complete denoising ($t = 089$), $t = 776$ reduces runtime from 71.29s to 33.42s while maintaining a competitive fusion score (SS 2.2431 vs. 2.3110). More generally, image quality improves rapidly in the early-to-middle stage and then saturates, whereas runtime continues to increase substantially.

Timestep	**Time(s)** ↓	**SS** ↑	**CLIP-FlanT5** ↑			**Qwen2.5-VL** ↑			**LLaVA-1.5** ↑		
			P_g	P_{single}	P_{Fuse}	P_g	P_{single}	P_{Fuse}	P_g	P_{single}	P_{Fuse}
999	**9.75**	1.5418	0.6439	0.4416	0.5067	0.5925	0.2607	0.4357	0.5633	0.3341	0.3313
971	13.66	1.7090	0.6377	0.4854	0.5359	0.6361	0.3451	0.4906	0.5603	0.3773	0.3407
943	16.60	1.7534	0.6623	0.5162	0.5742	0.6483	0.3895	0.5467	0.5724	0.4017	0.3650
913	20.23	2.0037	0.6890	0.5782	0.6369	0.6981	0.5212	0.7057	0.5987	0.4617	0.4271
881	23.51	2.0270	0.6825	0.5814	0.6466	0.7374	0.5825	0.7691	0.5938	0.4559	0.4338
848	26.62	2.1644	0.6991	0.6021	0.6628	0.7914	0.6455	0.8245	0.6096	0.4955	0.4579
813	29.97	2.2041	0.6983	0.5927	0.6603	0.7916	0.6313	0.8209	0.6030	0.4851	0.4533
776†	33.42	2.2431	0.6958	0.5967	0.6577	0.7967	0.6341	0.8219	0.6011	0.4916	0.4634
736	36.09	2.2193	0.6828	0.5818	0.6439	0.8027	0.5914	0.8031	0.5969	0.4826	0.4512
695	39.21	2.2234	0.7015	0.6052	0.6763	0.7849	0.6131	0.8350	0.6067	0.4788	0.4635
651	43.21	2.2618	0.6972	**0.6139**	0.6754	0.7924	0.6293	0.8317	0.6076	**0.4984**	**0.4674**
604	46.18	2.2705	**0.7051**	0.6062	0.6764	0.7992	0.6084	0.8470	0.6060	0.4825	0.4654
554	50.71	2.2660	0.7021	0.6039	0.6690	0.8054	0.6155	0.8337	0.6050	0.4723	0.4595
501	53.12	2.2471	0.6987	0.5935	0.6641	0.8045	0.6126	0.8330	0.5987	0.4663	0.4445
444	56.62	2.2633	0.7025	0.6003	0.6623	0.8106	0.6187	0.8470	0.6092	0.4695	0.4586
383	59.60	2.2606	0.6972	0.5900	0.6603	**0.8168**	0.6087	0.8273	0.6035	0.4733	0.4523
318	63.17	2.2389	0.6933	0.5904	0.6632	0.8122	0.6018	0.8220	0.6033	0.4749	0.4505
247	67.58	2.2954	0.6933	0.6050	0.6786	0.8127	0.6349	0.8424	**0.6106**	0.4864	0.4648
171	69.29	2.2845	0.7023	0.6094	**0.6804**	0.8162	0.6281	0.8501	0.6088	0.4866	0.4619
089	71.29	**2.3110**	0.6961	0.6119	0.6735	0.8146	**0.6457**	**0.8553**	0.6065	0.4966	0.4635

RFigure 1. **Preliminary three-image fusion results.** We extend ITMDiff to three inputs using a simple sequential strategy: two images are first fused into an intermediate hybrid, which is then fused with a third image. The results suggest that the framework can be extended beyond pairwise fusion, while also showing more noticeable information loss and imbalance as the number of inputs increases.